

# Locality Sensitive Hashing

Rebecca C. Steorts

October 12, 2017

## 1 Locality Sensitive Hashing (LSH)

A *hash function* maps objects to integers such that dissimilar objects are mapped far apart. LSH uses special hash functions to ensure that similar objects are put close to each other (with high probability). By applying several locality-sensitive hashes to a record, one goes from a high-dimensional object to a low-dimensional signature, with a guarantee that similar records have nearby signatures. The low-dimensional signature vectors can then be divided into bins or blocks, with a high probability that all records mapped to the same bin are similar, and a hope that not too many similar records fall into different bins. If the number of records per bin is small, we can treat the bins as blocks for purposes of record linkage. The number of records per bin depends on the exact binning procedure and its tuning parameters (Kim and Lee 2010; Rajaraman and Ullman 2012; Christen 2012). Unlike conventional blocking, LSH uses all the fields of a record and can be adjusted to ensure that blocks are manageably small.<sup>1</sup> By design, records falling within the same block are similar to each other, so only linking records within blocks can lead to dramatic speed-ups while imposing only a small cost in false negative errors. (One does fewer comparisons, and those comparisons are more likely to lead to links.)

While these are all desirable features, LSH is in several ways not yet fully developed as a blocking technique for entity resolution. One is that “similarity” must be made computationally precise, and not every similarity measure is preserved by a (known) family of hash functions. Moreover, different similarity measures may be appropriate for different kinds of data in different applications. Second, entity resolution should come not from ranking similarity scores (as in Liang et al. 2014), but from a statistical

---

<sup>1</sup>The last point needs some care, since simulation studies show that *sometimes* bins can contain many records, leading to minimal computational savings.

30 model. We develop a variety of LSH blocking algorithms which are designed  
31 to work well with real-world examples of entity resolution and to meet the  
32 needs of my statistical models.

## 33 1.1 Minwise hashing

34 Minwise hashing is a procedure for comparing the similarity of two sets  
35 (records in this context can be thought of as sets). For a large dataset like  
36 the Syrian one we are considering, minwise hashing is optimal because of its  
37 speed in subdividing records into blocks. The implementation of minwise  
38 hashing has a number of steps.

### 39 First

40 We shingle each record. Shingling is a way of splitting apart a record  
41 into smaller subsets, at some specified splitting point. For example,  
42 if we have the record ["Peter", "Pittsburgh"], and we want to use  
43 a shingling of size  $k = 3$ , we form the shingles ["Pet", "erP", "itt",  
44 "sbu", "rgh"]. To do this, we combine the entire record into a single  
45 string, and then subdivide it into parts of three.

46 Shingling is done to each record in the dataset, which forms a large set  
47 of shingles. We also want to keep track of which records are associated  
48 with which shingles. So if we have a second record, ["Steve", "Pitts-  
49 burgh"], then we know that the shingles "itt", "sbu", and "argh" are  
50 in common between the two records.

### 51 Second

52 We form a characteristic matrix from the shingles. A characteristic  
53 matrix is an indicator matrix, where the rows of the matrix correspond  
54 to records in the dataset (so there is a row for each record) and the  
55 columns to the shingles formed from the records (so there is a column  
56 for each shingle). The elements of the matrix are a binary (1, 0)  
57 decision, where a 1 is present if the record contains the corresponding  
58 shingle, and a 0 if not. The characteristic matrix is very sparse (i.e.,  
59 it contains mainly 0s), as most records do not contain the majority of  
60 all possible shingles.

61 Table 1.1 is an example of a characteristic matrix, with four shingles  
62 and five records.

### 63 Third

64 We permute the rows of the characteristic matrix to form a permuted

Row	$S_1$	$S_2$	$S_3$	$S_4$
1	0	0	1	1
2	1	0	0	1
3	1	1	0	0
4	0	0	0	1
5	1	0	0	1

Table 1: A sample characteristic matrix, with four shingles and five possible records. As the number of shingles grows, the matrix becomes increasingly sparse, as most records do not contain most shingles.

matrix. If we have five rows, 1-5, then possible permutations of the rows are 12345, 15234, and 54321 (there are many others). The permuted matrix, then, is simply a reordering of the original characteristic matrix, with the rows swapped in some arrangement. Figure 1.1 shows the characteristic matrix converted to a permuted matrix by a given permutation.

We repeat the permutation step for several iterations to obtain multiple permuted matrices.

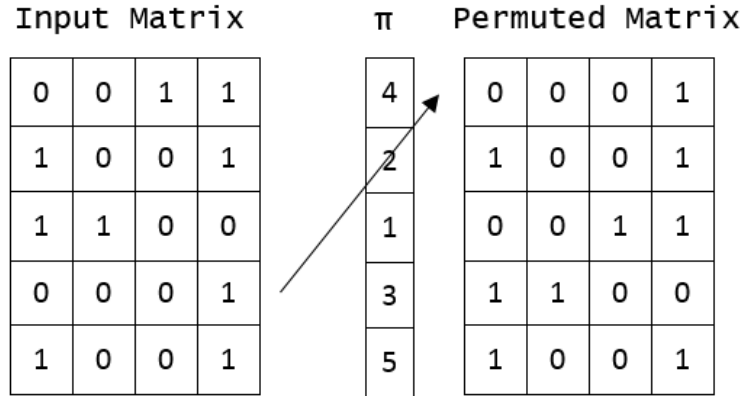


Figure 1: Permuted matrix from the characteristic one. The  $\pi$  vector is the specified permutation.

In practice, the implementation of minhashing is slightly different from these permutations, as permuting the matrices a large number of times

75 is computationally-prohibitive. To make up for this, a series of hash  
76 functions is generated, where the hash function assign values to each  
77 record, in exactly the same way as the permutations do. The hash  
78 function,  $(r + 1)\%5$ , for example, takes in a record index (1-5), and  
79 assigns the new indices as the permuted matrix. This results in the  
80 same outcomes as the one above. An example of these hash functions  
81 with the characteristic matrix is given in table 1.1.

Row	$S_1$	$S_2$	$S_3$	$S_4$	$(r+1) \% 5$	$(3r+2) \% 5$
0	0	0	1	1	1	2
1	1	0	0	1	2	0
2	1	1	0	0	3	3
3	0	0	0	1	4	1
4	1	0	0	1	0	4

Table 2: Characteristic matrix with four sets, two hash functions, and a generated binary matrix. The hash values are found by inputting the corresponding row index to the hash function.

## 82 Fourth

83 We compute the signature matrix. The signature matrix is a hashing  
84 of values from the permuted one. The signature has a row for the  
85 number of permutations calculated, and a column corresponding with  
86 the columns of the permuted matrix. We iterate over each column of  
87 the permuted matrix, and populate the signature matrix, row-wise,  
88 with the row index from the first 1 value found in the column. The  
89 row index inputted to the signature matrix is the original row index  
90 the row was associated with, rather than the new index value. The  
91 signature matrix for table 1.1 is in table 1.1.

	$S_1$	$S_2$	$S_3$	$S_4$
$h_1$	2	3	1	4

Table 3: Signature matrix formed from table 1.1. The value for each element is the row index in which the first 1 is found in the permuted matrix.

## 92 Fifth

93 Based on the results of the signature matrix, we observe the pairwise  
 94 similarity of two records. For the number of hash functions, there is  
 95 an interesting relationship that we use between the columns for any  
 96 given set of the signature matrix and a Jaccard Similarity measure.

The Jaccard Similarity, which is a measure of similarity, is fundamental in this approach. For two sets S and T, the Jaccard Similarity is given by

$$\frac{|S \cap T|}{|S \cup T|}$$

97 which is the intersection of the two sets over their union. It provides  
 98 a sense of how much two records agree in their fields, over their total  
 99 number of variables.

100 The relationship between the random permutations of the character-  
 101 istic matrix and the Jaccard Similarity is:

$$Pr\{min[h(A)] = min[h(B)]\} = \frac{|A \cap B|}{|A \cup B|}$$

102 The equation means that the probability that the minimum values of  
 103 the given hash function, in this case  $h$ , is the same for sets A and  
 104 B is equivalent to the Jaccard Similarity, especially as the number of  
 105 record comparisons increases. The output of this formula, in terms of  
 106 the signature matrix, can be seen in figure 1.1.




Figure 2: Two sets S and T with Jaccard similarity 2/5. The two sets share two elements in common, and there are five elements in total.

107 We use this relationship to calculate the similarity between any two  
 108 records. We look down each column, and compare it to any other col-

109        umn: the number of agreements over the total number of combinations  
110        is equal to Jaccard measure.

111        Of course, in practice, we would need to repeat the above for many per-  
112        mutations  $\pi$  such that we avoid collisions, or mapping the different records  
113        into the same bin.

114        Notice that we have walked through a very elementary view of locality  
115        sensitive hashing, and we still have the limitation that we are having to do  
116        all-to-all comparisons. How do we avoid this?

## 117    2    LSH for Minhash Signatures

118        See Section 3.4.1 in Mining for Massive Datasets and work through the  
119        exercises (some of these we did in class). See Example 3.11 and Exercise  
120        3.4.1 and 3.4.2.

## 121    3    Further reading

122        For further reading about hashing and it's use in practice, please refer to  
123        <https://arxiv.org/abs/1710.02690>.

## 124    References

- 125        CHRISTEN, P. (2012). *Data Matching: Concepts and Techniques for Record Link-*  
126        *age, Entity Resolution, and Duplicate Detection*. Springer.
- 127        KIM, H.-S. and LEE, D. (2010). Harra: fast iterative hashed record linkage for  
128        large-scale data collections. In *Proceedings of the 13th International Conference*  
129        *on Extending Database Technology*. ACM, 525–536.
- 130        LIANG, H., WANG, Y., CHRISTEN, P. and GAYLER, R. (2014). Noise-tolerant ap-  
131        proximate blocking for dynamic real-time entity resolution. In *Eighteenth Pacific-*  
132        *Asia Conference on Knowledge Discovery and Data Mining (PAKDD14)* (Z.-H.  
133        Zhou, A. L. P. Chen, V. S. Tseng and B. H. Tu, eds.). Springer, New York, NY,  
134        forthcoming.
- 135        RAJARAMAN, A. and ULLMAN, J. D. (2012). *Mining of massive datasets*. Cam-  
136        bridge University Press.