

An Introduction to Information Retrieval, Part II

STA 325, Supplemental Material

Announcements

- ▶ Exam I grades, summaries, etc have been posted
- ▶ Preparing for datathon

Datathon

- ▶ You can find some preliminary resources for datathon here:
`https://github.com/resteorts/data-mine/tree/master/datathon-2018`
- ▶ Above you will find a template for your presentation and your final write up.

Groups for Datathon

You'll find group leaders below for datathon as well as your groups.

```
groupLeaders <- c("Donheiser, Julia ", "Wohl, Julia", "Rajan, Hari",  
  "Donahoe, Grace", "Chen, Carolyn", "Holland, Ethan",  
  "Ben Alexander")
```

```
Donheiser <- c("Agarwal, Naman", "Deng, Ziqi", "Lin, Vincent")  
Wohl <- c("Gao, An", "Lyu, Cheng", "Xu, Yingzhi")  
Rajan <- c("Metz, Audreya", "Giraud, Paul", "Ma, Calvin")  
Donahoe <- c("Marsalek, Oskar", "Gu, Yunxiang", "Shmueli, Sara")  
Chen <- c("San-Martin, Jose", "Hsiao, Man-Lin", "Tan, Michael")  
Holland <- c("Taylor, Moriah", "Huang, Howard", "Wang, Harry")  
Alexander <- c("Malone, Alexis", "Jiang, Debra", "Xiao, Claire")
```

Datathon

1. Recommendations: Work on upcoming homeworks with your group members.
2. Today: work on the in class assignment with your group (part of homework 5 so you can better prepare for the datathon).
3. Group leader should be in charge of delegating tasks to other members of the group.
4. Feel free to ask myself/Ben questions when you work on in class assignments to prepare for the datathon.

Grading for Datathon

1. Everyone must be present for all of Datathon.
2. All code and documents must be kept on a github repository. Please add Ben and I as collaborators. Please keep these as private until the group presentations occur.
3. Your grade will be calculated on what you come up with during the Datathon event (from start to finish).
4. You'll be allowed to make some changes, which I'll outline below after Datathon, but these should be related to making your code reproducible, cleaning up your report/slides, fixing any visualizations that you may have had trouble with during the Datathon.
5. Please submit your original reports from Datathon and your final reports to Sakai by the due date on the Syllabus. (Wednesday, November 7, 10 PM)

Questions

1. What if I have a really neat idea for improving my analysis after the Datathon? (This is great! Mention this in the future work section/conclusions of your report. You're welcome to spend extra time trying things out and placing these in the appendix. Your grade will be calculated solely on what you did the day of Datathon).
2. What if our analysis didn't work?

Make sure to state this in your report and what your analysis didn't work. What would you recommend if you had more time? If your analysis did not work much at all, but sure to elaborate on this part more.

3. How should I start revising my report?

Feel free to chat with me about this with your group and I'm happy to make suggestions.

Grading for Datathon

- ▶ Datathon is worth 30 percent of your course grade

You'll be graded on datathon based upon two components: a written report (20 percent) and an oral presentation (10 percent)

Remark: If a group member is not participating in any component of datathon, the report, or oral presentation, the group leader should email me. In the event, your group leader is not being active, any group member should contact me to let me know ASAP.

Written report (20 percent)

Your grade will be based upon 5 components:

1. Reproducibility of code and document (20 points)
2. Exploratory analysis provided on data set and description of data set (20 points)
3. Creativity, originality of proposed methods (20 points)
4. Analysis applied to the data set and conclusions made (visualizations, evaluations) (20 points)
5. Oral presentation (clarity, ease of reading) (20 points)

Oral presentation (10 percent)

1. 5 percent of your grade will be based upon your slide submission
 2. 5 percent of your grade will be based upon your oral presentation to the class
- ▶ Both of the written and oral presentations are due to Sakai by your group on Wednesday November 7, 10 PM (Sakai).

Note: all components of the written and oral presentation must be reproducible. Failure to do so will result in a 20 point penalty, so please check this in your group before submission.

- ▶ Your oral presentations to the class (6 minutes/group) will be on December 4 and 6.

Midterm course evaluations

At the end of the class today, if you would take the time to complete the midterm course evaluations, I would be very grateful. (You should have the link in your email to the survey).

Agenda

1. The exercise for today will integrate with homework 5.
2. We will work thorough parts of it so that you can work in your datathon teams.
3. The rest you will finish on your own, but feel free to meet up with your datathon team.
4. Our task today will be analyzing NYTimes articles about music and art and using our skills relating to information retrieval. We'll also be using our skills with functional programming. In essence, the goal of this homework is to put all of our knowledge thus far in the class to use!

Install packages

```
library("xtable")
```

```
## Warning: package 'xtable' was built under R version 3.4
```

Information retrieval applied to NY Times

Let's consider applying our information retrieval skills to a corpus of documents from the NY Times.

NY Times corpus

Let's start by seeing what kind of data we have. To do this, we're going to use some pre-written functions located in `ir.R` that you are welcome to use.

Loading scripts

```
source("scripts/ir.R")
```


Reading in the data

```
library(XML)
music <- read.directory("nyt_corpus/music")
art <- read.directory("nyt_corpus/art")
```

Let's examine the data

```
length(music)
```

```
## [1] 45
```

```
length(art)
```

```
## [1] 57
```

We have 45 music articles and 57 art articles from the NYTimes.

Let's examine the data

1. What command would you use to extract the 37th word of story number 1595645 in art? (That word is “experiencing”.)
2. Give a command to count the number of times the word “the” appears in that story.

(Try this on your own and don't look at the solution.)

Solution

```
which(dir("nyt_corpus/art")== "1595645.xml")
```

```
## [1] 48
```

The 37th word is

```
art[[48]][37]
```

```
## [1] "experiencing"
```

Solution (continued)

One way to see how often “the” appears in the story above is using the following command:

```
sum(art[[48]] == "the")
```

```
## [1] 103
```

Bag of words

3. Give the commands you would use to construct a bag-of-words data-frame from the document vectors for the art and music stories.

(Try this on your own and don't look at the solution.)

Solution

`lapply` is one of the most useful functions in R: it takes two arguments, a data structure (vector, list, array, ...) and a function, and returns a list which it gets by applying the function to each element of the data structure.

```
art.bow <- lapply(art,table)
music.bow <- lapply(music,table)
is.list(music.bow)
```

```
## [1] TRUE
```

```
length(music.bow)
```

```
## [1] 45
```

Solution (continued)

```
nyt.frame <- make.BoW.frame(c(art.bow,music.bow))  
dim(nyt.frame)
```

```
## [1] 102 4431
```

The number of rows should equal the total number of stories, and $45+57 = 102$.

Distance matrices

Create distance matrices from this data frame for (a) the straight Euclidean distance, (b) the distance with word-count normalization and (c) the distance with vector-length scaling, and then for all three again with inverse-document-frequency weighting. Be sure to give the commands that you use. ¹

¹Note: There are six different distance commands that can be found in the function `distances`. Some of these were covered in class and some of these are new ones that you have not seen before that you will explore during the assignment.

Partial solution

Let's just look at seeing how to tackle this problem.

```
dist.plain = distances(nyt.frame)
dim(dist.plain)
```

```
## [1] 102 102
```

```
round(dist.plain[1:5,1:5],2)
```

```
##      [,1]  [,2]  [,3]  [,4]  [,5]
## [1,]  0.00 163.82 116.99 129.27 132.92
## [2,] 163.82  0.00  87.26  87.42  78.60
## [3,] 116.99  87.26  0.00  68.26  78.89
## [4,] 129.27  87.42  68.26  0.00  80.46
## [5,] 132.92  78.60  78.89  80.46  0.00
```

```
dist.wordcount = distances(div.by.sum(nyt.frame))
dist.euclen = distances(div.by.euc.length(nyt.frame))
```

On your own

Now, do the idf weights on your own.

On your own

```
nyt.frame.idf = idf.weight(nyt.frame)
dist.idf.plain = distances(nyt.frame.idf)
dist.idf.wordcount = distances(div.by.sum(nyt.frame.idf))
dist.idf.euclen = distances(div.by.euc.length(nyt.frame.idf))
```

Average distances between stories

For each of the six different difference measures in the function `distances`, what is the average distance between stories in the same category and between stories in different categories?

Solution

There are multiple ways to do this.

The simplest is to realize that, in this case, the first 57 stories are all art, and the last 45 are all music.

- ▶ So if d is a distance matrix, the within-category entries are $d[1:57, 1:57]$ and $d[58:102, 58:102]$, and the between-category entries are $d[1:57, 58:102]$.
- ▶ Then $\text{mean}(c(d[1:57, 1:57], d[58:102, 58:102]))$ would give the average distance between stories in the same category, similarly $\text{mean}(d[58:102, 1:57])$ for the between-category average.

Solution

The `outer()` function takes three arguments: two data-structures and another function. (Here the function is `!=`, which I put in quotes so that R realizes I'm naming a function, and not asking it to evaluate an expression.)

It returns a matrix which it gets from applying the function to each pair of components from its first two arguments.

Here those first two arguments are vectors of length 102, so what it gives back is a 102×102 matrix, where `are.different[i,j]` shows whether `class.label[i] != class.label[j]`. In other words, it's TRUE if documents `i` and `j` belong to different classes.

Solution

```
class.labels = c(rep("art",57),rep("music",45))  
head(class.labels)
```

```
## [1] "art" "art" "art" "art" "art" "art"
```

```
are.different = outer(class.labels,class.labels,"!=")  
head(are.different)
```

```
##      [,1] [,2] [,3] [,4] [,5] [,6] [,7] [,8] [,9] [,10] [,11] [,12] [,13] [,14] [,15] [,16] [,17] [,18] [,19] [,20] [,21] [,22]  
## [1,] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE  
## [2,] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE  
## [3,] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE  
## [4,] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE  
## [5,] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE  
## [6,] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE  
##      [,12] [,13] [,14] [,15] [,16] [,17] [,18] [,19] [,20] [,21] [,22] [,23] [,24] [,25] [,26] [,27] [,28] [,29] [,30] [,31]  
## [1,] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE  
## [2,] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE  
## [3,] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE  
## [4,] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE  
## [5,] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE  
## [6,] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```


Solution

And a logical array picks out elements from another array:
`mean(d[are.different])` is the average distance between classes.
To average within classes, `mean(d[!are.different])` Not only does this work if the classes are intermingled (we just have to get the `class.labels` vector right), we can also use this to not include the distance from a document to itself in the within-class average:

With-in class differences (with self)

```
# calculate with-in class differences, with self  
are.same = !are.different  
diag(are.same) = TRUE  
mean(dist.plain[are.same])
```

```
## [1] 76.43675
```

```
mean(dist.wordcount[are.same])
```

```
## [1] 0.101356
```

```
mean(dist.euclen[are.same])
```

```
## [1] 0.7293624
```

```
mean(dist.idf.plain[are.same])
```

```
## [1] 87.93946
```

```
mean(dist.idf.wordcount[are.same])
```

With-in class differences (without self)

```
# calculate with-in class differences, without self  
are.same = !are.different  
diag(are.same) = FALSE  
mean(dist.plain[are.same])
```

```
## [1] 77.94421
```

```
mean(dist.wordcount[are.same])
```

```
## [1] 0.1033549
```

```
mean(dist.euclen[are.same])
```

```
## [1] 0.7437465
```

```
mean(dist.idf.plain[are.same])
```

```
## [1] 89.67376
```

```
mean(dist.idf.wordcount[are.same])
```

Between category averages

```
# calculate between category averages  
mean(dist.plain[58:102,1:57])
```

```
## [1] 78.06805
```

```
mean(dist.wordcount[58:102,1:57])
```

```
## [1] 0.1069678
```

```
mean(dist.euclen[58:102,1:57])
```

```
## [1] 0.7670132
```

```
mean(dist.idf.plain[58:102,1:57])
```

```
## [1] 88.25976
```

```
mean(dist.idf.wordcount[58:102,1:57])
```

```
## [1] 0.1322932
```

Output into an xtable

```
are.same = !are.different
diag(are.same) = TRUE
wselfwdif<-c(mean(dist.plain[are.same]),
             mean(dist.wordcount[are.same]),
             mean(dist.euclen[are.same]),
             mean(dist.idf.plain[are.same]),
             mean(dist.idf.wordcount[are.same]),
             mean(dist.idf.euclen[are.same]))
are.same = !are.different
diag(are.same) = FALSE
wselfwodif<-c(mean(dist.plain[are.same]),
             mean(dist.wordcount[are.same]),
             mean(dist.euclen[are.same]),
             mean(dist.idf.plain[are.same]),
             mean(dist.idf.wordcount[are.same]),
             mean(dist.idf.euclen[are.same]))
betweendif<-c(mean(dist.plain[58:102,1:57]),
             mean(dist.wordcount[58:102,1:57]),
             mean(dist.euclen[58:102,1:57]),
             mean(dist.idf.plain[58:102,1:57]),
             mean(dist.idf.wordcount[58:102,1:57]),
             mean(dist.idf.euclen[58:102,1:57]))
mat<-cbind(wselfwdif, wselfwodif, betweendif)
rownames(mat)<-c("Plain {without IDF}",
               "Sum-normed {without IDF}", "Length-normed
               {without IDF}", "Plain {with IDF}", "Sum-normed
               {with IDF}", "Length-normed {with IDF}")
colnames(mat)<-c("Within-class, with self", "Within-class, without
               self", "Between-class")
```

Print xtable

```
print(xtable(mat))
```

```
## % latex table generated in R 3.4.1 by xtable 1.8-3 package
## % Sun Oct 28 11:13:44 2018
## \begin{table}[ht]
## \centering
## \begin{tabular}{rrrr}
## \hline
## & Within-class, with self & Within-class, without & sel
## \hline
## Plain \{without IDF\} & 76.44 & 77.94 & 78.07 \\\
## Sum-normed \{without IDF\} & 0.10 & 0.10 & 0.11 \\\
## Length-normed
## \{without IDF\} & 0.73 & 0.74 & 0.77 \\\
## Plain \{with IDF\} & 87.94 & 89.67 & 88.26 \\\
## Sum-normed
## \{with IDF\} & 0.13 & 0.13 & 0.13 \\\
## Length-normed \{with IDF\} & 1.35 & 1.37 & 1.38 \\\
## \hline
## \end{tabular}
## \end{table}
```

Multidimensional scaling plots

Create multidimensional scaling plots for the different distances, and describe what you see. (Include the code you used, the plots, and explanations for the code).

Solution

Recall the point of multi-dimensional scaling is such that we can take our distance matrix and represent this in lower dimensions, namely two, so that we can visualize it.


Solution

```
# we run classical MDS on our distance matrices  
head(cmdscale(as.matrix(dist.plain)))
```


```
##           [,1]      [,2]  
## [1,] 142.885340  3.068330  
## [2,]  -6.909667  8.437926  
## [3,]  53.400617 -4.257171  
## [4,]  47.720333 -3.421029  
## [5,]  47.205833  5.192395  
## [6,]  35.295081  3.055806
```

Solution


No normalization, no weighting




Word-count normalization, no weighting




Euclidean length normalization, no weighting




No normalization, IDF weighting



Word-count normalization, IDF weighting




Euclidean length normalization, IDF weighting



Top row, without IDF. Bottom row, with IDF. Left column, un-normalized vectors. Middle column, normalized by word-count. Right column, normalized by Euclidean length. Red circles are art,

Solution

Euclidean length normalization, IDF weighting



Only with both IDF weights and Euclidean-length normalization do we get reasonable separation of the two categories.

Summary

- ▶ We have completed a first analysis of the NYTimes articles in this class (and assignment).
- ▶ In homework 5 you'll re-inforce this knowledge on your own and complete a few other exercises.